# Hacktoberfest 2026 trades PRs for events

Hacktoberfest 2026 no longer counts pull requests or merge requests toward rewards, replacing the twelve-year "open N PRs, earn a T-shirt" model with month-long, MLH- and DEV-run events and challenges under an "AI belongs to everyone" mission funded by DigitalOcean as presenting partner. In its place, online participants earn a mailed sticker pack at 3, 10, and 15 (or 17, per MLH) virtual stickers among 24 collectible, participation-driven stickers, while 300+ one-day Fests across 60+ countries, cash-prize DEV Challenges, Global Hack Week (Oct 9-15), and livestreams anchor the program ((FAQ) (https://hacktoberfest.com/questions/), (MLH blog) (https://blog.mlh.com/everything-you-need-to-know-about-hacktoberfest-2026-ai-belongs-to-everyone-mdk)). The event is free and open to anyone aged 13+ (subject to country and Fest-specific rules), yet measurable outcomes remain sparse: the most reliable public figures as of mid-October 2026 are the MLH events API count (531 Fests across 71 countries) and 669,307 open hacktoberfest issues (~517,632 unclaimed) on aggregator sites, not an official participation total ((MLH events API) (https://hacktoberfest-api.mlh.com/api/events), (goodfirstissue.org) (https://www.goodfirstissue.org/labels/hacktoberfest)).

## PR rewards vanish as AI spam overwhelms maintainers

The defining change in 2026 is the abolition of contribution counting. "Pull requests and merge requests will no longer count toward Hacktoberfest rewards," because low-effort spam PRs became trivial to generate, and organizers say they are "listening to maintainer feedback" ((Hacktoberfest FAQ) (https://hacktoberfest.com/questions/), (Home) (https://hacktoberfest.com/)). The MLH announcement is blunt: maintainers began to burn out under the volume of low-effort, box-checking pull requests, and AI tools made generating those PRs "trivial" while creating "no value" ((MLH blog, Aug 19, 2026) (https://blog.mlh.com/hacktoberfest-2026-ai-belongs-to-everyone-3jl8)). This justification rests on a documented history: 2020 saw 621,104 October PRs with 9,598 flagged spam/invalid and 172,599 sent to repos that never opted in, prompting introduction of the hacktoberfest topic opt-in, hacktoberfest-accepted labels, and a 7-day review window ((Developer Relations case study) (https://developerrelations.com/case-studies/overcoming-hacktoberfest-spam/)). By 2025, the eligible-PR target had climbed from 4 to 6 to quality-gate the T-shirt, itself capped at the first 10,000 "super contributors", and the OpenSSF-attributed estimate that ~35% of Hacktoberfest-labeled PRs were closed without merging illustrates the chronic signal-to-noise problem organizers now cite ((openresource.dev) (https://openresource.dev/articles/hacktoberfest/), (tamiz.pro analysis) (https://tamiz.pro/insights/hacktoberfest-2026-quality-first-open-source-habits)). Historical participation had already declined nearly two-thirds, from 169,886 registered participants in 2020 to 56,768 in 2025, with eligible PRs falling from 483,127 (2019) to ~87,929 (2025) at roughly 1.55 PRs per participant ((tvc-16.science) (https://tvc-16.science/rip-hacktoberfest.html)). Whether PR-removal itself caused new pain is unverified; the best evidence

is mixed, and tvc-16 argues the participation collapse tracks tech labor-market contraction more than event mechanics ((tvc-16.science) (https://tvc-16.science/rip-hacktoberfest.html)).

## In-person Fests and online challenges replace contributions

With no PR leaderboard, legitimacy flows through calendars, Discord, and in-person rooms. Hack Day and Meetup Fests are one-day events (Hack Days run 3-12 hours with demos and teams) where hosts receive MLH funding/reimbursement while Meetups run unfunded, and participants gain certificates, swag while supplies last, and possibility of Hack Day prizes (Hacktoberfest T-shirts, sensor-based Arduino projects at select events) ((In-person overview) (https://hacktoberfest.com/in-person/), (Host a Fest) (https://hacktoberfest.com/host/), (FAQ) (https://hacktoberfest.com/questions/)). The MLH events API showed 531 events across 71 countries as of mid-October 2026-led by India (224) and the US (101), with India 12 and Bengaluru 12 events individually prominent-meaning geographic access is highly uneven and regional capacity can be as low as hand-picked 50 participants at some India events ((MLH events API) (https://hacktoberfest-api.mlh.com/api/events), (spiketoday.com) (https://spiketoday.com/hacktoberfest-rules/), events.mlh API descriptions]). Online, Global Hack Week: Hacktoberfest (Oct 9-15, 2026) delivered daily livestreamed workshops, design/technical/registration/social challenges, and Discord mini-events with GHW points ([(GHW events) (https://ghw.mlh.com/events/open-source)), while DEV runs five back-to-back challenges-a Weekend Challenge (submission deadline Oct 4, 11:59 PM PDT) plus four weekly themes (Week 1 Oct 5-11, Week 2 Oct 12-18, Week 3 Oct 19-25, Week 4 Oct 26-31) each offering Overall $250 + DEV++ + badge, Featured $200, category $100, and completion badges-all built on the single prompt "Build something with open-source AI at its core" ((DEV HF26) (https://dev.to/challenges/hf26), (DEV Week 1) (https://dev.to/challenges/hacktoberfest-week1-2026-10-05)). Free partner credits (Tinker, Render, Backboard.io, .tech domains) add a sponsor-centric flavor to the build incentives ((DEV Week 1) (https://dev.to/challenges/hacktoberfest-week1-2026-10-05)).

## Tagged issue backlogs keep contributions flowing without rewards

The official pivot did not end repo-side October work, but it reclassified it. GoodFirstIssue.org indexed 669,307 open hacktoberfest issues on Oct 16, 2026, of which roughly 517,632 remain unclaimed, with an average annotated difficulty of 3.7/5 and 11% of time estimates at one day or less; ~9,738 issues opened in the prior seven days show the backlog is actively replenished ((goodfirstissue.org) (https://www.goodfirstissue.org/labels/hacktoberfest)). The GitHub hacktoberfest topic covers 148,820 public repositories (with 25,097 carrying hacktoberfest-accepted) (GitHub topics (https://github.com/topics/hacktoberfest)), and TrendForge’s Oct 11 weekly #hacktoberfest report surfaced 46 trending projects with star deltas, including Apollo-11, Termux, Kestra, and Bubble Tea (TrendForge (https://trendforge.devlive.org/report/topic/hacktoberfest/weekly/2026-10-11)). Maintainer communities (Owncast, TunaOS, KubeStellar, GDG chapters, OpenCodeChicago, Python Philippines/RELX, React Hyderabad) publish opt-in guidance, and a DEV maintainer hub frames

"self-contained issues you can actually finish in a sitting" for AI agent tooling (TunaOS hacktoberfest 2026 (https://tunaos.org/docs/tunaos/hacktoberfest-2026), Chawla DEV post (https://dev.to/tisha/hacktoberfest-2026-help-us-govern-and-reproduce-ai-agents-59n1)). Common beginner deliverables are documentation, unit tests, i18n string wrapping, magic-number extraction, and checksum/SBOM docs, accepted via claim-before-work rules, DCO sign-off (git commit -s), linked issues (Fixes #123), focused diffs, and CI tests (TunaOS (https://tunaos.org/docs/tunaos/hacktoberfest-2026), Scrimba beginner guide (https://scrimbaguide.tech/blog/hacktoberfest-2026-first-pull-request/)). Aggregators recommend issue-hunter queries like is:issue is:open label:"hacktoberfest" label:"good first issue" no:assignee plus -linked:pr, since visible discoverability of unclaimed issues-not issue creation-is now the critical Curation metric (boomspot guide (https://www.boomspot.com/how-to-find-good-first-issues-for-hacktoberfest-2026)). The pivot thus shifted open-source October effort from counting to matching well-scoped, claimable work, while the event itself no longer endorses any specific repo or label scheme (hacktoberfest.com/resources/maintainers 404 (https://hacktoberfest.com/resources/maintainers)).

## Controversy and impact debates split open-source voices

The 2026 pivot is contested on pipelines, access, and corporate influence. Sponsors include Tiger Data, Snowflake, MongoDB, Gauge, Solana, Render, GitHub, Sentry, Backboard.io, IBM, ElevenLabs, Paper Compute, Entire, Prior Labs, Google Cloud/Gemma, Qualcomm, Arduino, Mastra, Temporal, TLDR, Thinking Machines, .tech, and SerpApi, creating a narrative of an event increasingly "built by partners" (llms-full.txt (https://hacktoberfest.com/llms-full.txt)).
Critics on Reddit complain that losing PR tracking "means the death of Hacktoberfest as a phenomenon," citing loss of review-based feedback and motivation (r/hacktoberfest comment (https://www.reddit.com/r/hacktoberfest/comments/1vsuffw/hacktoberfest_2026_ai_belongs_to_everyone/)), and India contributors note that Fest availability lags India’s size ("very less fests compared to size of India") with disappointment over the format change (r/developersIndia (https://www.reddit.com/r/developersIndia/comments/1w913n6/are_there_really_no_hacktoberfest_2026_events/)). Historical retainers such as no T-shirt promise online (only in-person supplies and a Completionist raffle of a T-shirt or an Arduino Uno Q) added grievance, since 2025 had briefly reintroduced T-shirts capped at 10,000 before the 2026 cut (FAQ (https://hacktoberfest.com/questions/), (tvc-16.science) (https://tvc-16.science/rip-hacktoberfest.html)). Organizers still cite educational wins: Jess Lee reported 500+ Fests in 60+ countries and 3,000 projects submitted in the first week, including a Pakistan Hack Day hospital platform and a multimodal accessibility tool (DEV context (https://dev.to/dailycontext/one-weekend-3000-open-source-ai-projects-3g65)). The causal debate over impact is unresolved: Hacktoberfest historically powered hundreds of thousands of first contributions (freeCodeCamp grew a quiz site from 694 to 1,171 questions in one October; Meilisearch hired contributor Bruno Casali ten days after his first contribution in 2021), while challenges like ServiceNow’s saw 1,375 PRs (+89%) in 2024, but 2026’s new model deliberately de-targets those pipelines (openresource.dev (https://openresource.dev/articles/hacktoberfest/), ServiceNow Community (https://www.servicenow.com/community/developer-advocate-blog/hacktoberfest-2024-review/ba-p/3099284)). Open-resource skeptics worry the event trades

first-PR pathways to local events and cash-prize contests, and that AI-slop pressures (already compared to 2020’s "Shitoberfest") persist year-round, with one maintainer describing a "first assumption on receiving a large PR is that it’s AI-generated," eroding willingness to review (martinsos.com (https://martinsos.com/posts/contributing-in-2026), Publicasta (https://publicasta.com/open_source_radar/ai_slop_pull_requests_open_source_maintainers_2026_08_28?lang=en)). Accessibility provision is delegated to hosts; explicit channels exist via "Contact Host" for accommodation/equipment requests, beginner Meetups require no AI project experience, and a community (Cameroon, month-long, multi-hub) agenda explicitly includes accessibility and African-language topics while discouraging mass-generated or unexplainable AI contributions (hacktoberfest.cm (https://hacktoberfest.cm/), (In-person) (https://hacktoberfest.com/in-person/)). No Hacktoberfest-specific 2026 demographic data have been published, and the event remains in progress, so verdicts on inclusion, pipeline outcomes, or spam reduction are unverifiable from official data (tvc-16.science (https://tvc-16.science/rip-hacktoberfest.html)).

## Best practices now favor bounded scope and owned contributions

Experienced maintainers and contributors converge on operational rigor rather than PR volume. Recommended maintainer practices include: conduct contribution gates in CI (tests at all appropriate levels, no warnings, docs updates when behavior changes), keep PRs scoped to roughly under ~400 lines, publish a CONTRIBUTING.md, issue templates with required fields, and difficulty tiers, audit CI/CD and consider OpenSSF Scorecards, plus saved close messages ("does not meet our contribution bar") and clear AI-disclosure policies that require reproduction details, versions, and exploitability for security issues (tamiz.pro/dev.to (https://tamiz.pro/insights/hacktoberfest-2026-quality-first-open-source-habits), Publicasta (https://publicasta.com/open_source_radar/ai_slop_pull_requests_open_source_maintainers_2026_08_28?lang=en), Open resource retrospective (https://openresource.dev/articles/hacktoberfest/)). Historically effective spam counters-hacktoberfest topic opt-in, hacktoberfest-accepted acceptance labels, 7-day review windows, spam labels, reporting tools, advisory council-remain the playbook even though 2026 deleted PR counting entirely, and GitHub’s 2026 PR open limits and interaction limits for non-write-access users provide infrastructure-level throttling (Developer Relations (https://developerrelations.com/case-studies/overcoming-hacktoberfest-spam/), openresource.dev (https://openresource.dev/articles/hacktoberfest/)). Contributor-side advice centers on claiming issues explicitly and waiting for acknowledgment, respecting scope, writing tests, linking issues, and not submitting cosmetic-only sweeps; disclosing AI use and owning each contribution (use the project first, write without AI "smell" when possible) builds the trust that maintainers treat as scarce (Help Wanted guide (https://helpwanted.dev/blog/hacktoberfest-2026), Chawla DEV post (https://dev.to/tisha/hacktoberfest-2026-help-us-govern-and-reproduce-ai-agents-59n1), martinsos.com (https://martinsos.com/posts/contributing-in-2026)). One cloud maintainer’s 2025 experience shows the failure mode when AI-assisted PRs are approved in bulk: one PR "fixed" an already-resolved, unclosed issue likely fed to Cursor, and review stretched beyond Nov 15, reinforcing that claimed scope and pre-vetted issue subsets reduce waste (Cloudinary DEV post (https://dev.to/cloudinary/hacktoberfest-as-a-first-time-maintainer-2pkn)).

## Conclusion

Hacktoberfest 2026 marks a deliberate separation of community participation from PR volume accounting, betting that in-person mentorship and AI build-and-ship challenges can sustain momentum while AI makes synthetic contributions cheap. The tradeoff is explicit: upstream projects lose the October "first contribution" flood that historically fed long-term pipelines, but gain protection from volume-driven review collapse, and contributors gain claimable, bounded work and local onboarding but lose reward-activated contribution feedback. Whether the new model reinvigorates participation or marks managed decline depends on metrics organizers have not yet published; decisive evidence will come from post-October retrospectives, Fest attendance data, DEV Challenge submission quality, and maintainer-reported review load, not from announcement enthusiasm. The lasting lesson is that reward design is the central lever of open-source participation contests, and that 2026 optimizes for AI literacy and local networks rather than upstream diff counts.